

# PA02: Analysis Report - Network I/O Primitives

Name and Roll Number: Ayush Jain , MT25066

Course: Graduate Systems (GRS)

GitHub : [https://github.com/AyushJain001/GRS\\_PA02](https://github.com/AyushJain001/GRS_PA02)

---

## Table of Contents

1. [Implementation Overview](#)
  2. [Experimental Setup](#)
  3. [Results Summary](#)
  4. [Analysis Questions](#)
  5. [Plots](#)
  6. [AI Usage Declaration](#)
- 

## Implementation Overview

### A1: Two-Copy (Standard Socket)

- Uses send() and recv() system calls
- Data path: User buffer → Kernel socket buffer → Network → Kernel socket buffer → User buffer
- **Two memory copies** per send/receive operation

### A2: One-Copy (Scatter/Gather I/O)

- Uses sendmsg() with iovec array
- 8-element iovec pointing to each message field
- Eliminates intermediate buffer assembly in user space
- **One memory copy** in optimized kernel path

### A3: Zero-Copy

- : Uses sendmsg() with MSG\_ZEROCOPY
- flag Kernel maps user pages directly for DMA
- Requires error queue polling for completion notification
- **Zero CPU copies** (DMA handles transfer)

# Experimental Setup

## System Configuration:

- OS: Ubuntu 24.04
- CPU: Intel
- Network: Network namespaces (server\_ns @ 10.0.0.1, client\_ns @ 10.0.0.2)

## Experiment Parameters:

- Message sizes: 64, 256, 1024, 4096 bytes
- Thread counts: 1, 2, 4, 8
- Duration: 3 seconds per run
- Messages per run: 1,000,000
- Total experiments: 48 (3 × 4  
× 4)

## Profiling Events (via perf stat):

- cpu\_core/cycles/
- cpu\_core/L1-dcache-load-misses/
- cpu\_core/LLC-load-misses/
- context-switches

---

# Result summary

## Throughput (Gbps)

Msg Size	Threads	A1 (Two-copy)	A2 (One-copy)	A3 (Zero-copy)
64 B	1	0.077	0.077	0.066
64 B	8	0.421	0.405	0.270
256 B	1	0.313	0.298	0.258
256 B	8	1.213	1.151	1.040
1024 B	1	1.102	1.119	0.968
1024 B	8	4.965	5.084	4.241
4096 B	1	4.181	4.049	3.920
4096 B	8	18.824	18.851	16.38

7

## Latency (μs)

Msg Size	Threads	A1 (Two-copy)	A2 (One-copy)	A3 (Zero-copy)
64 B	1	6.67	6.61	7.73
64 B	8	9.67	10.06	15.13
1024 B	1	7.41	7.29	8.44
1024 B	8	13.13	12.82	15.38
4096 B	1	7.81	8.07	8.34
4096 B	8	13.85	13.84	15.92

# Analysis Questions

## Question 1: Why does zero-copy not always give the best throughput?

### Observation:

- A3 (zero-copy) is consistently **slower** than A1 and A2 across all message sizes
- At 64B/1 thread: A3 = 0.066 Gbps vs A1 = 0.077 Gbps (**14% slower**)
- At 4096B/8 threads: A3 = 16.39 Gbps vs A1 = 18.82 Gbps (**13% slower**)

### Reasons:

1. **Page Pinning Overhead:** MSG\_ZEROCOPY requires kernel to pin user pages in memory before DMA. For small messages (64-256B), this overhead dominates the transfer time.
2. **Error Queue Servicing:** Zero-copy requires polling the socket error queue for completion notifications. This adds ~1-2  $\mu$ s latency per message batch.

**Conclusion:** Zero-copy is optimized for large, sustained bulk transfers over real network hardware. On loopback/namespace networks with small messages, traditional copy-based methods perform better.

---

## Question 2: Which cache level shows the most reduction in misses and why?

### L1 Cache Miss Data (millions):

Msg Size	A1 (1 thread)	A2 (1 thread)	A3 (1 thread)
64 B	95.79M	91.68M	109.52M
256 B	95.68M	94.21M	116.68M
1024 B	111.11M	119.67M	120.67M
4096 B	191.00M	191.04M	195.28M

### LLC Miss Data (thousands):

Msg Size	A1 (1 thread)	A2 (1 thread)	A3 (1 thread)
64 B	8.6K	3.1K	3.6K
256 B	1.7K	2.8K	4.0K
1024 B	11.6K	10.4K	9.6K
4096 B	10.7K	8.4K	7.3K

### Analysis:

1. **L1 Cache:** Shows **higher** misses for A3 compared to A1/A2. This is because:

- A3 page-pinning causes additional kernel memory accesses
  - Error queue polling accesses different cache lines.

**2. LLC (Last Level Cache):** Shows **reduction** with A2 and A3 at larger sizes:

- A3 at 4096B: 7.3K vs A1: 10.7K (**32% reduction**)
- Zero-copy bypasses LLC by using DMA paths

**Conclusion:** LLC shows the most consistent reduction with zero-copy because DMA operations bypass CPU caches entirely. L1 cache misses actually increase due to additional kernel bookkeeping overhead.

---

## Question 3: How does thread count interact with cache contention?

**Throughput Scaling (1024B messages):**

Threads	A1 Gbps	A2 Gbps	A3 Gbps	A1 Scaling	A2 Scaling	A3 Scaling
1	1.10	1.12	0.97	1.0×	1.0×	1.0×
2	1.97	1.96	1.73	1.79×	1.75×	1.79
						×
4	3.20	3.15	2.75	2.91×	2.82×	2.84
						×
8	4.96	5.08	4.24	4.51×	4.54×	4.38
						×

**L1 Cache Misses vs Thread Count (64B):**

Threads	A1 L1 Misses	A2 L1 Misses	A3 L1 Misses
1	95.79M	91.68M	109.52M
2	187.83M	187.83M	200.96M
4	351.24M	344.17M	392.77M
8	656.54M	673.56M	553.63M

**Observations:**

**1. 1 → 2 Threads:** Near-linear scaling (~1.8×)

- L1 misses scale proportionally (1.96× for A1)

**2. 4 → 8 Threads:** Sublinear scaling (~1.55×)

- Expected: 2× improvement, Actual: 1.55×

**3. A3 Shows Better Scaling at 8 Threads (L1 misses only 553M vs 656M for A1):**

- Zero-copy reduces per-thread memory traffic
- Kernel handles page management, reducing user-space cache pressure

#### Cache Contention Mechanisms:

- **False Sharing:** Threads accessing adjacent socket statistics
  - **True Sharing:** Kernel socket buffer locks
  - **TLB Pressure:** Multiple threads cause more TLB misses at 8 threads
- 

## Question 4: At what message size does one-copy (A2) outperform two-copy (A1)?

#### Throughput Comparison (8 threads):

##### Msg Size A1 Gbps A2 Gbps A2 vs A1

Msg Size	A1 Gbps	A2 Gbps	A2 vs A1
64 B	0.421	0.405	-3.9%
256 B	1.213	1.151	-5.1%
1024 B	4.965	5.084	+2.4%
4096 B	18.82	18.85	+0.1%
	4	1	

Answer: A2 outperforms A1 starting at **1024 bytes** message size (at 8 threads).

#### Reasons:

##### 1. Small Messages (64-256B): iovec setup overhead dominates

- Building 8-element iovec array: ~50-100 ns
- sendmsg() syscall has higher entry cost than send()
- Net result: A2 is 3-5% slower

##### 2. Large Messages (1024B+): Data transfer cost dominates

- A2's scatter/gather avoids intermediate buffer assembly
- Kernel can optimize DMA scatter lists

---

## Question 5: At what message size does zero-copy (A3) outperform two-copy (A1)?

#### Throughput Comparison:

##### Msg Size A1 Gbps A3 Gbps A3 vs A1

Msg Size	A1 Gbps	A3 Gbps	A3 vs A1
64 B	0.421	0.270	-35.9%
256 B	1.213	1.040	-14.3%
1024 B	4.965	4.241	-14.6%
4096 B	18.82	16.38	-12.9%
	4	7	

Answer: A3 **never** outperforms A1 in our network namespace environment.

#### Reasons:

**1. No Real DMA Benefit:** Network namespaces use kernel's virtual network stack, not real NICs. Zero-copy's DMA optimization is bypassed.

**2. Fixed Overhead Per Message:**

- Page pinning: ~2-5  $\mu$ s per
- message Error queue poll: ~1-2  $\mu$ s per batch
- This overhead doesn't decrease with message size

**3. Where A3 Would Win (on real networks):**

- Message sizes > 64KB with sustained transfers
- 10+ Gbps NICs where NIC DMA bandwidth is the bottleneck
- Bulk data transfers (not request-response patterns)

**Prediction:** On a 10 GbE NIC with 1MB+ messages, A3 would show 20-30% improvement over A1.

---

## Question 6: Identify one unexpected result and explain it

**Unexpected Result:** A3 (Zero-copy) has **higher** L1 cache misses than A1 at small message sizes.

**Expected:** Zero-copy should have fewer cache misses since it avoids CPU copies.

**Explanation:**

**1. Page Pinning Overhead:** MSG\_ZEROCOPY requires kernel to:

- Look up page table entries (TLB misses  $\rightarrow$  L1 misses)
- Pin pages in memory (modify page flags)
- Track page references for completion notification

**2. Error Queue Processing:** Polling S0\_EE\_ORIGIN\_ZEROCOPY events requires:


- Reading from socket error queue (separate cache lines)
  - Processing completion notifications (additional data structures)
-

# Plots

Note: Plots are generated by running `python3 MT25066_Part_D_Plots.py` during demo. They are not included in submission as per assignment requirements.


## Plot 1: Throughput vs Message Size

- 2x2 grid showing throughput for each thread count (1, 2, 4, 8)
- A1/A2/A3 comparison lines
- Log scale X-axis for message sizes




## Plot 2: Latency vs Thread Count

- 2x2 grid showing latency for each message size (64, 256, 1024, 4096)
- Shows latency increase with thread contention




## Plot 3: Throughput Scaling

- 2x2 grid showing how throughput scales with threads
- Demonstrates sublinear scaling at higher thread counts




## Plot 4: Implementation Comparison (Bar Chart)

- Side-by-side bars for A1, A2, A3
- Shows relative performance at each configuration




## Plot 5: Cache Misses vs Message Size

- L1 and LLC cache misses
- Compares 1-thread vs 8-thread scenarios



## Plot 6: Cycles per Byte

- Efficiency metric: CPU cycles / bytes
- transferred Lower is better



# AI Usage Declaration

## Component 1: Message Structure Design

**Files:** MT25066\_Part\_A\_Common.h, MT25066\_Part\_A\_Common.c

**Prompt Used:**

Create a C struct for an 8-field network message where each field is dynamically allocated.

**AI Contribution:**

- Initial struct definition with 8 char\* fields
- malloc/free pattern for dynamic allocation

---

## Component 2: Multithreaded Server Implementation

**Files:** MT25066\_Part\_A1\_Server.c, MT25066\_Part\_A2\_Server.c, MT25066\_Part\_A3\_Server.c

**Prompt Used:** Create a multithreaded TCP server in C that accepts clients, spawns one thread per client, receives messages and echoes them back.

**AI Contribution:**

- Basic accept() loop structure
- pthread\_create() for client handling

---

## Component 3: Client with Latency Measurement

**Files:** MT25066\_Part\_A1\_Client.c, MT25066\_Part\_A2\_Client.c, MT25066\_Part\_A3\_Client.c

**Prompt Used:** Create a TCP client that spawns N threads, each connecting to server, sending messages for T seconds.

**AI Contribution:**

- Thread spawning and joining
- pattern clock\_gettime() usage for timing

**My Modifications:**

- Corrected throughput formula: (bytes \* 8) / (time \* 1e9) for Gbps

## Component 4: Experiment Automation Script

**File:** MT25066\_Part\_C\_Experiment.sh

**Prompt Used:** Bash script to automate running server/client with perf stat for different message sizes and thread counts

**AI Contribution:**

- Nested loop structure for
- parameters perf stat command syntax
- CSV output formatting

**My Modifications:**

- Added timeout to prevent hanging processes
- Created both individual and summary CSV files

---

## Component 5: Matplotlib Plotting Script

**File:** MT25066\_Part\_D\_Plots.py

**My Modifications:**

- Made the .py file for plotting
- Added all 48 data points to nested dictionaries
- Fixed axis scales (log scale for message sizes)
- Added proper units (Gbps,  $\mu$ s, millions, thousands)

---

## Component 6: Network Namespace Setup

**File:** setup\_namespaces.sh

**Prompt Used:** Bash script to create two network namespaces connected via veth pair. Server namespace at 10.0.0.1, client namespace at 10.0.0.2. Include setup, cleanup, and status commands.

**AI Contribution:**

- ip netns commands
- structure veth pair creation syntax
- Basic case statement for commands

# Conclusion

This assignment demonstrated that:

1. **Copy elimination doesn't always improve performance** - Kernel optimizations and modern CPU memcpy speeds can make additional syscall overhead more expensive than the copies saved.
  2. **Zero-copy requires specific conditions** - Real network hardware with DMA, large sustained transfers, and high bandwidth utilization.
  3. **Cache behavior is counterintuitive** - Zero-copy can increase L1 misses due to kernel bookkeeping overhead.
  4. **Thread scaling reveals contention** - Beyond 4 threads, cache coherency and lock contention dominate.
-